


# SmartX Labs for Computer Systems

Cluster &  
Analytics Lab  
(2018, Spring)  
NetCS Lab



# History and Contributor of Cluster Lab

## (2017. 05. 20.)


Version	Updated Date	Updated Contents	Contributor
-	2015/10	(구) Analytics Lab 작성	송지원
v1	2016/04	Cluster Lab 초안 작성	김승룡
v2	2016/05	Cluster Lab 수정	송지원
v3r3	2016/05/28	Cluster Lab 2차 수정 (내용 수정 및 추가)	송지원
v4r1	2016/05/30	Cluster Lab 3차 수정 (피드백 반영)	송지원
v5	2016/06/01	HDFS를 옵션으로 변경, 기타 문제 수정	송지원
v6r1	2016/06/03	실습자 검수 후 수정	송지원, 윤희범, 남택호
v6r2	2016/06/29	HDFS 설치 과정 등 수정	송지원
v6r3	2016/06/30	Zeppelin 독립 실행 모드 설명 추가	송지원
0.6.4	2016/07/04	Zeppelin 설치 방법 누락된 부분 추가	송지원
0.6.5	2017/05/20	Ubuntu 16.04, Spark 2.1.1-Hadoop-2.7, Zeppelin 0.7.1, Hadoop 2.8.0 대응 업데이트	강문중
0.6.6	2018/05/30	소프트웨어 버전 업데이트, 강의 시나리오 반영 하여 그림 변경, 이론 보강	권진철

# CSLab: Cluster & Analytics Lab.

## - Goal

SETUP to run data processing and visualization

- with Mesos, Spark, Zeppelin, (HDFS)



# Apache Mesos

## - Concept



MESOS

### What is Mesos?

**Apache Mesos** is an open-source project to manage **computer clusters**.

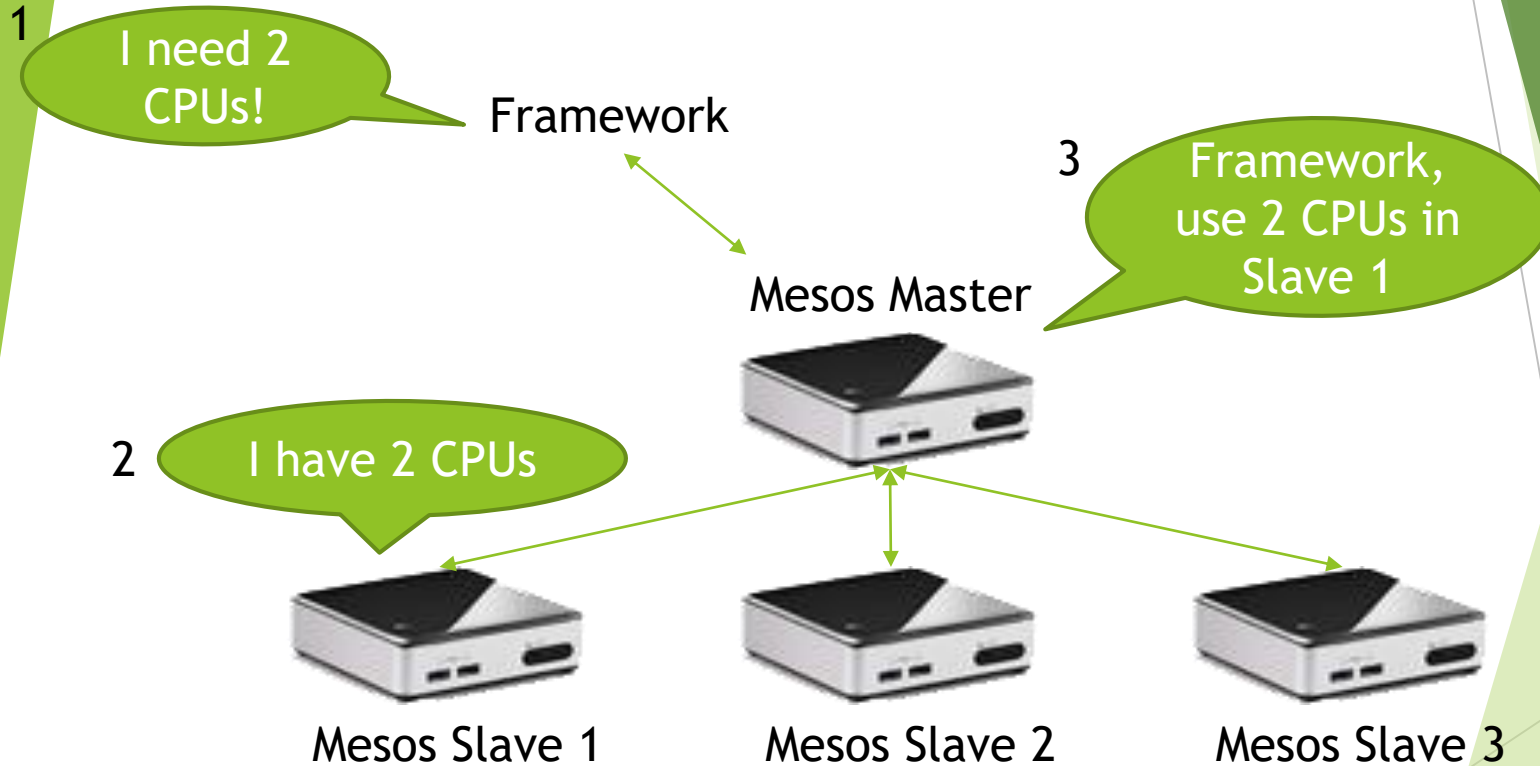
Mesos abstracts CPU, memory, storage, and other compute resources away from machines (physical or virtual), enabling fault-tolerant and elastic distributed systems to easily be built and run effectively.

Mesos is built using the same principles as the Linux kernel, only at a different level of abstraction. The Mesos kernel runs on every machine and provides applications (e.g., Hadoop, Spark, Kafka, Elastic Search) with API's for resource management and scheduling across entire datacenter and cloud environments.

- Cloud as a single computer
- Share resources across the machines

# Apache Mesos

## - Architecture




# Apache Spark

## - Concept



Apache Spark™ is a fast and general engine for large-scale data processing.


- In-memory data processing framework: Fast!
- Easy to use, community fastly growing
- Libraries: SQL and DataFrame, Streaming, MLlib, GraphX
- Run on standalone or Mesos, Yarn, etc
- Scala, Java, Python



# Apache Zeppelin -Concept

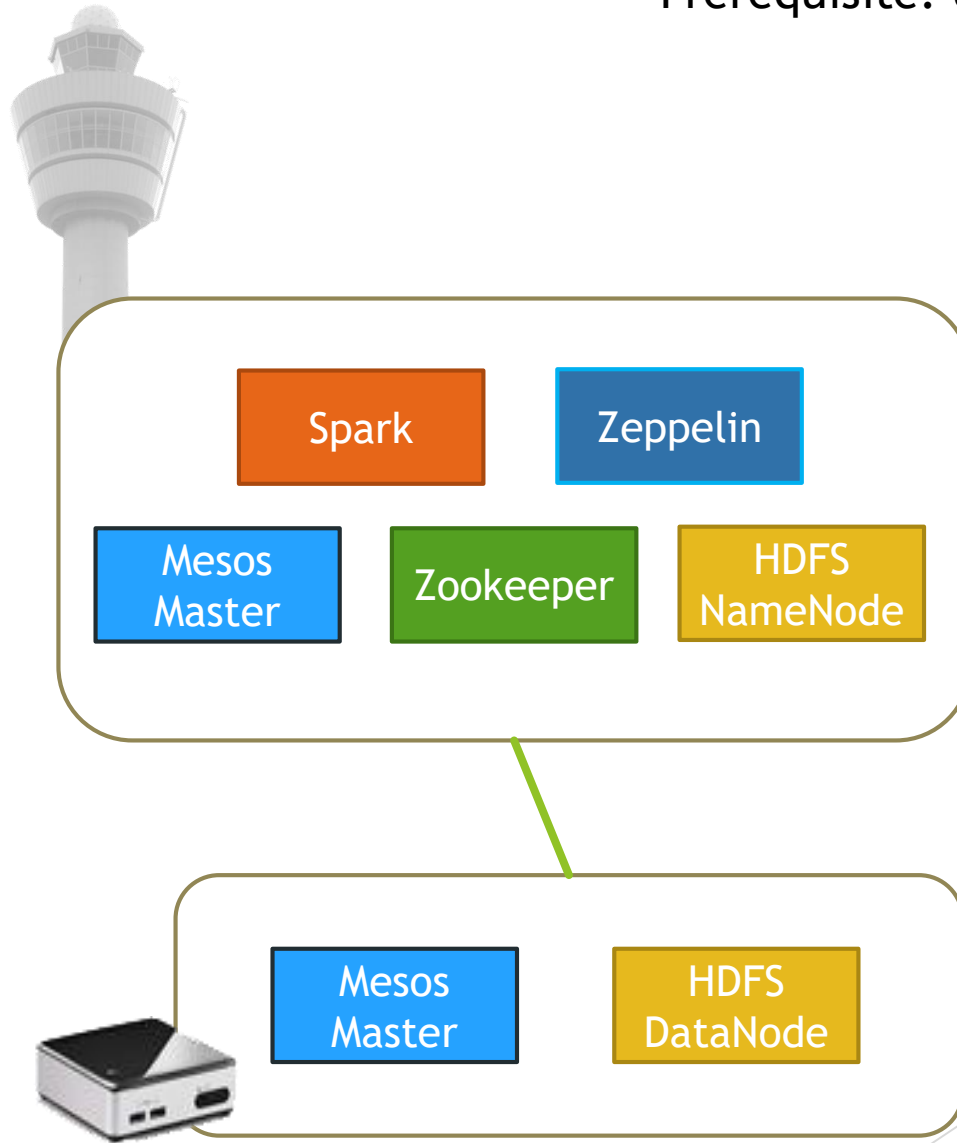
A web-based notebook that enables interactive data analytics.

Support Spark



# 0. Cluster Overview

Prerequisite: Ubuntu 16.04 - 64bit





# HDFS (optional)

## - Concept

### Hadoop Distributed FileSystem

- A distributed file system that provides high-throughput access to application data.

### Features


- **Fault tolerance** by detecting faults and applying quick, automatic recovery
- **Portability** across heterogeneous commodity hardware and operating systems
- **Scalability** to reliably store and process large amounts of data
- **Economy** by distributing data and processing across clusters of commodity personal computers
- **Efficiency** by distributing data and logic to process it in parallel on nodes where data is located
- **Reliability** by automatically maintaining multiple copies of data and automatically redeploying processing logic in the event of failures

# HDFS

## - Architecture

### <Master/Slave architecture>


- NameNode: A single node which manages the file system namespace and regulates access to files by clients.
- DataNode: DataNodes manage storage attached to the nodes that they run on.



# HDFS

## - Architecture

HDFS makes storages of separate machines in cluster into a single storage.



# 0. Preparation

## - Install Java and Mesos Dependencies



### Install JDK 8 and other Apache Mesos Dependencies

```
$ sudo apt update
$ sudo apt-get install -y openjdk-8-jdk
$ sudo apt-get -y install build-essential python-dev python-six python-virtualenv
libcurl4-nss-dev libsasl2-dev libsasl2-modules maven libapr1-dev libsvn-dev zlib1g-dev
iputils-ping
```

Do this for **all NUCs**.

# 0. Preparation

## - Configure hostnames



1. From **NUC 1** :

```
$ sudo hostname nuc01
```

2. From **NUC 2** :

```
$ sudo hostname nuc02
```

3. Edit /etc/hosts from **all NUCs** :

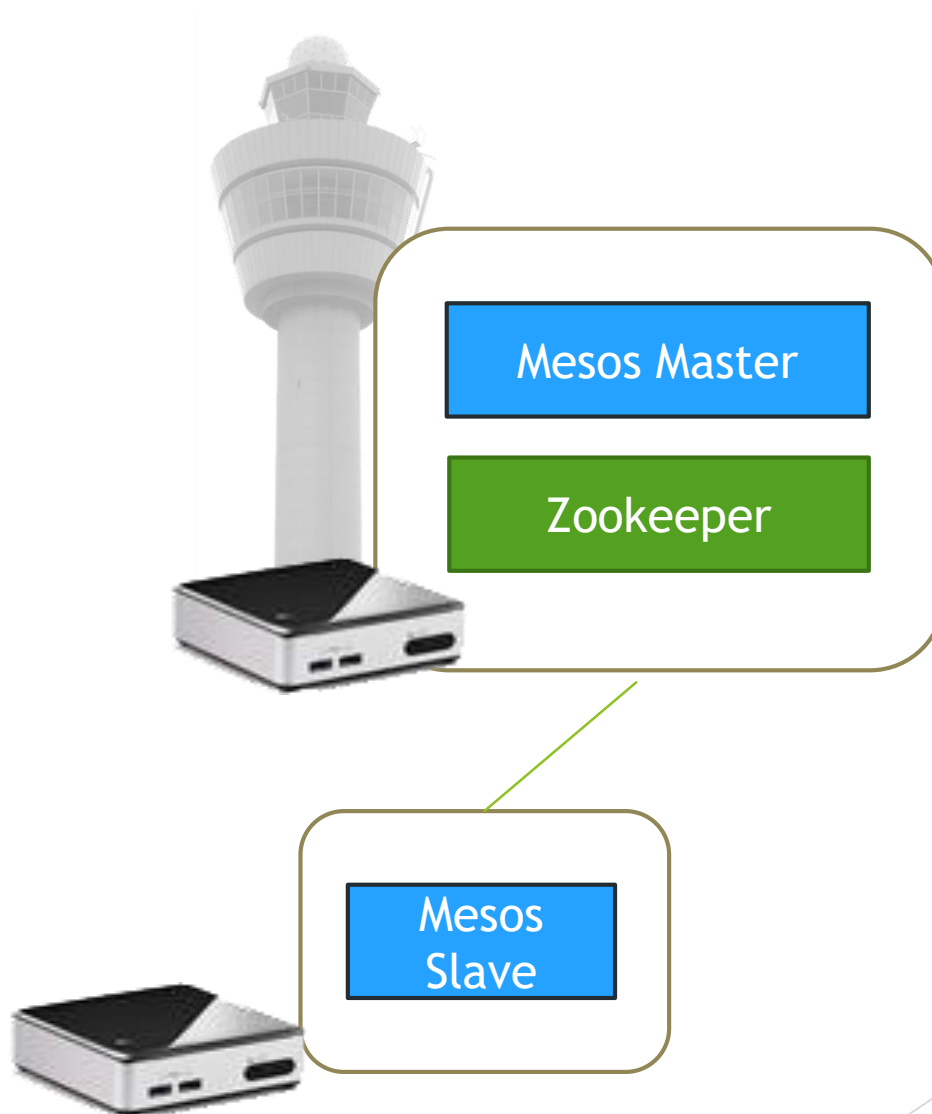
```
$ sudo vi /etc/hosts
```

4. Append the following context into /etc/hosts :

```
127.0.0.1      localhost
(IP Address of NUC 1)  nuc01
(IP Address of NUC 2)  nuc02
```

# 1. Apache Mesos

- Install



# 1. Apache Mesos

## - Installation Procedure

1. Add Mesosphere repository
2. Install **Mesos Master** on NUC 1
3. Install **Mesos Slave** on NUC 2
4. Check on Mesos Web UI

# 1. Apache Mesos

## - Install: Add Mesosphere repository

Add the repository to **all NUCs**.



```
$ sudo apt-key adv --keyserver keyserver.ubuntu.com --recv E56151BF
$ export DISTRO=$(lsb_release -is | tr '[:upper:]' '[:lower:]')
$ export CODENAME=$(lsb_release -cs)
```

To check you have correctly inputed, use the command below :

```
$ echo $DISTRO $CODENAME
```

Its result will be “ubuntu xenial”

```
$ echo "deb http://repos.mesosphere.io/${DISTRO} ${CODENAME} main" |
sudo tee /etc/apt/sources.list.d/mesosphere.list
$ sudo apt update
```



# 1. Apache Mesos

## - Install: Mesos Master on NUC 1



```
$ sudo apt -y install mesos
```

```
$ echo manual | sudo tee /etc/init/mesos-slave.override
```

```
$ echo 0.0.0.0 | sudo tee /etc/mesos-master/ip
```

```
$ echo nuc01 | sudo tee /etc/mesos-master/hostname
```

```
$ echo zk://localhost:2181/mesos | sudo tee /etc/mesos/zk
```

```
$ echo <NAME> | sudo tee /etc/mesos-master/cluster
```

```
$ echo 1 | sudo tee /etc/zookeeper/conf/myid
```

```
$ sudo systemctl restart zookeeper
```

```
$ sudo systemctl start mesos-master
```

<NAME>: anything you want

# 1. Apache Mesos

## - Install: Mesos Slave on NUC 2



```
$ sudo apt -y install mesos

$ echo manual | sudo tee /etc/init/mesos-master.override
$ echo 0.0.0.0 | sudo tee /etc/mesos-slave/ip
$ echo nuc02 | sudo tee /etc/mesos-slave/hostname
$ echo zk://<NUC1 IP>:2181/mesos | sudo tee /etc/mesos/zk
$ sudo cp /etc/mesos/zk /etc/mesos-slave/master
$ echo HADOOP_HOME=/usr/local/hadoop | sudo tee -a
  /etc/default/mesos-slave

$ sudo systemctl stop zookeeper
$ sudo systemctl start mesos-slave
```

HADOOP\_HOME will be  
needed later if you want  
to use Spark with HDFS.

# 1. Apache Mesos

## - Check on Mesos Web UI

In your web browser, go to

`http://<NUC 1 IP Address>:5050`



The screenshot shows the Apache Mesos Web UI for a cluster named 'nuc-cluster'. The top navigation bar includes links for Frameworks, Agents, Offers, and Maintenance. The current view is the 'Master' page, identified by the ID 'd0a2852a-3a16-41e6-a99e-8aaa357bf0a1'.

**Cluster Information:**

- Cluster: nuc-cluster
- Leader: nuc01:5050
- Version: 1.2.0
- Built: a month ago by ubuntu
- Started: 3 hours ago
- Elected: 3 hours ago

**Agents:**

State	Count
Activated	1
Deactivated	0
Unreachable	0

**Tasks:**

State	Count
Staging	0
Starting	0
Running	0
Unreachable	0
Killing	0
Finished	5
Killed	0
Failed	0
Lost	0
Orphan	0

**Active Tasks:**

Framework ID	Task ID	Task Name	State	Started	Host
No active tasks.					

**Unreachable Tasks:**

Framework ID	Task ID	Task Name	Started	Agent ID
No unreachable tasks.				

**Completed Tasks:**

Framework ID	Task ID	Task Name	State	Started	Stopped	Host
No completed tasks.						

**Orphan Tasks:**

Framework ID	Task ID	Task Name	State	Started	Stopped	Host
No orphan tasks.						

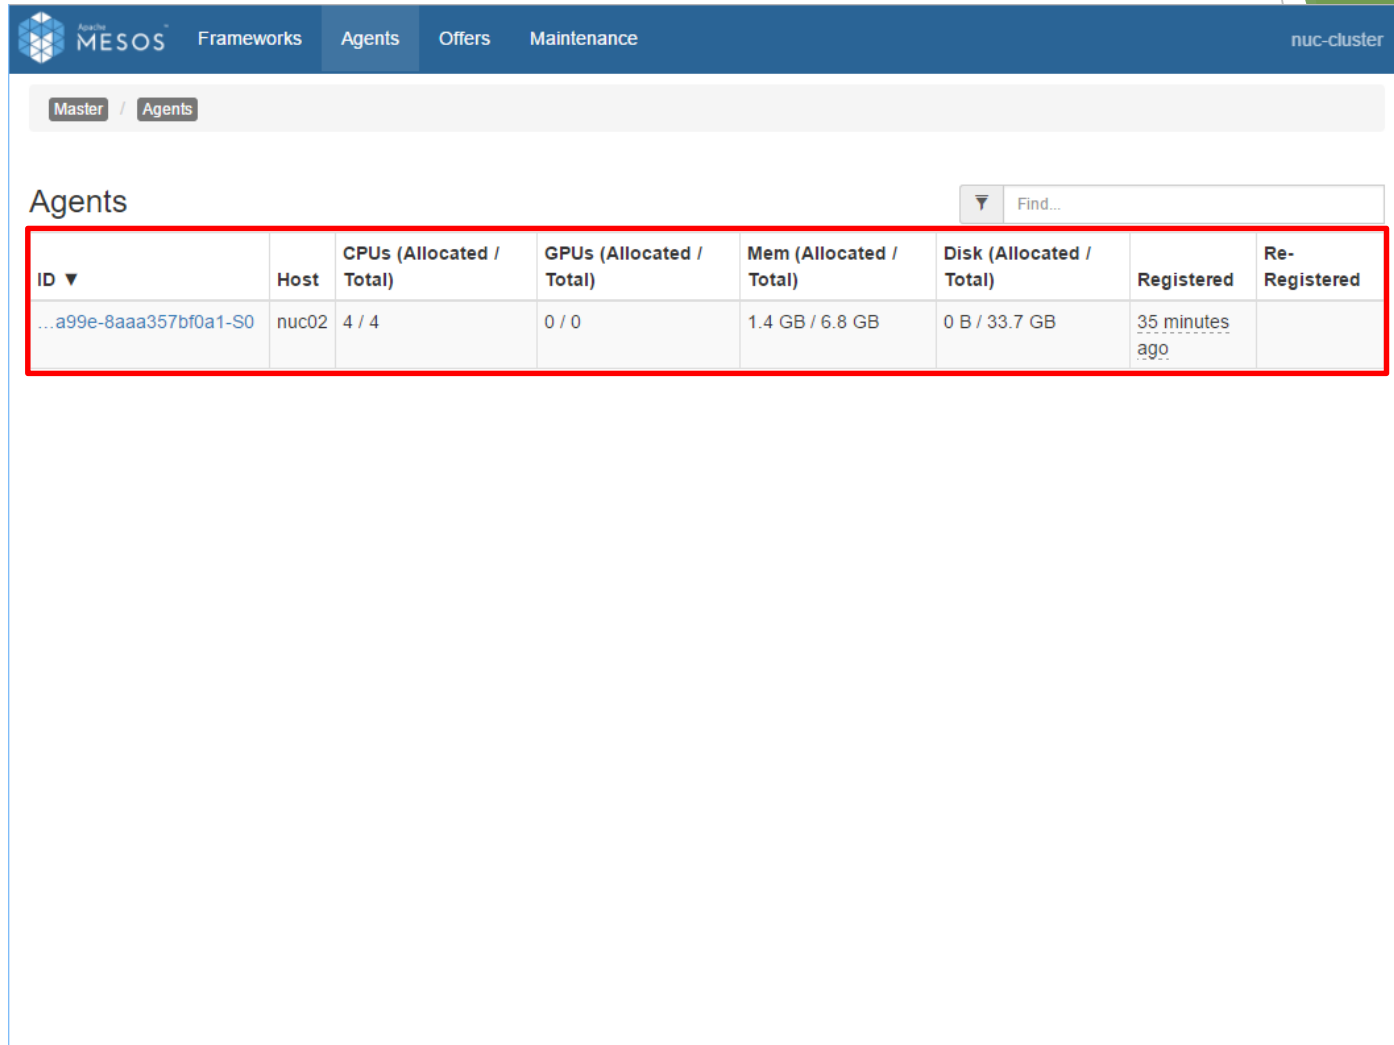
Check activated slaves and resources.

Note : In case of using the browser from other than those 2 NUCs, "hosts" file must be updated for the computer like the previous "Configure hostnames" step.

# 1. Apache Mesos

- Check on Mesos Web UI

Agents tab shows NUC 2 registered as a slave



The screenshot shows the Apache Mesos Web UI with the 'Agents' tab selected. The table below lists the registered agents. The first agent, 'nuc02', is highlighted with a red border.

ID ▼	Host	CPUs (Allocated / Total)	GPUs (Allocated / Total)	Mem (Allocated / Total)	Disk (Allocated / Total)	Registered	Re-Registered
...a99e-8aaa357bf0a1-S0	nuc02	4 / 4	0 / 0	1.4 GB / 6.8 GB	0 B / 33.7 GB	35 minutes ago	

# 2. Apache Spark

## - Installation Procedure

1. Install on NUC 1
2. Test on NUC 1

# 2. Apache Spark

## - Install

2. On **NUC 1**, Download and unarchive Spark, and configure spark-env.sh and spark-default.conf file.

```
$ cd ~
$ wget http://apache.mirror.cdnetworks.com/spark/spark-2.2.1/spark-2.2.1-bin-hadoop2.7.tgz
$ tar xzf spark-2.2.1-bin-hadoop2.7.tgz
$ cd spark-2.2.1-bin-hadoop2.7/conf
$ cp spark-env.sh.template spark-env.sh
$ cp spark-defaults.conf.template spark-defaults.conf
```

2.1. For **spark-env.sh**, append the following with vi command :

**Edit:**

- export SPARK\_LOCAL\_IP="**<NUC 1 IP Address>**"
- export MESOS\_NATIVE\_JAVA\_LIBRARY="/usr/local/lib/libmesos.so"
- export  
SPARK\_EXECUTOR\_URI="http://apache.mirror.cdnetworks.com/spark/spark-2.2.1/spark-2.2.1-bin-hadoop2.7.tgz"

2.2. For **spark-defaults.conf**, append the following with vi command :

**Edit:**

- spark.master mesos://**<NUC 1 IP Address>**:5050



# 2. Apache Spark

## - Test on NUC 1

# Start PySpark

```
$ cd ..  
$ bin/pyspark
```

# See if PySpark is running well

```
> data = range(1, 10001)  
> distData = sc.parallelize(data)  
> distData.filter(lambda x: x < 10).collect()
```

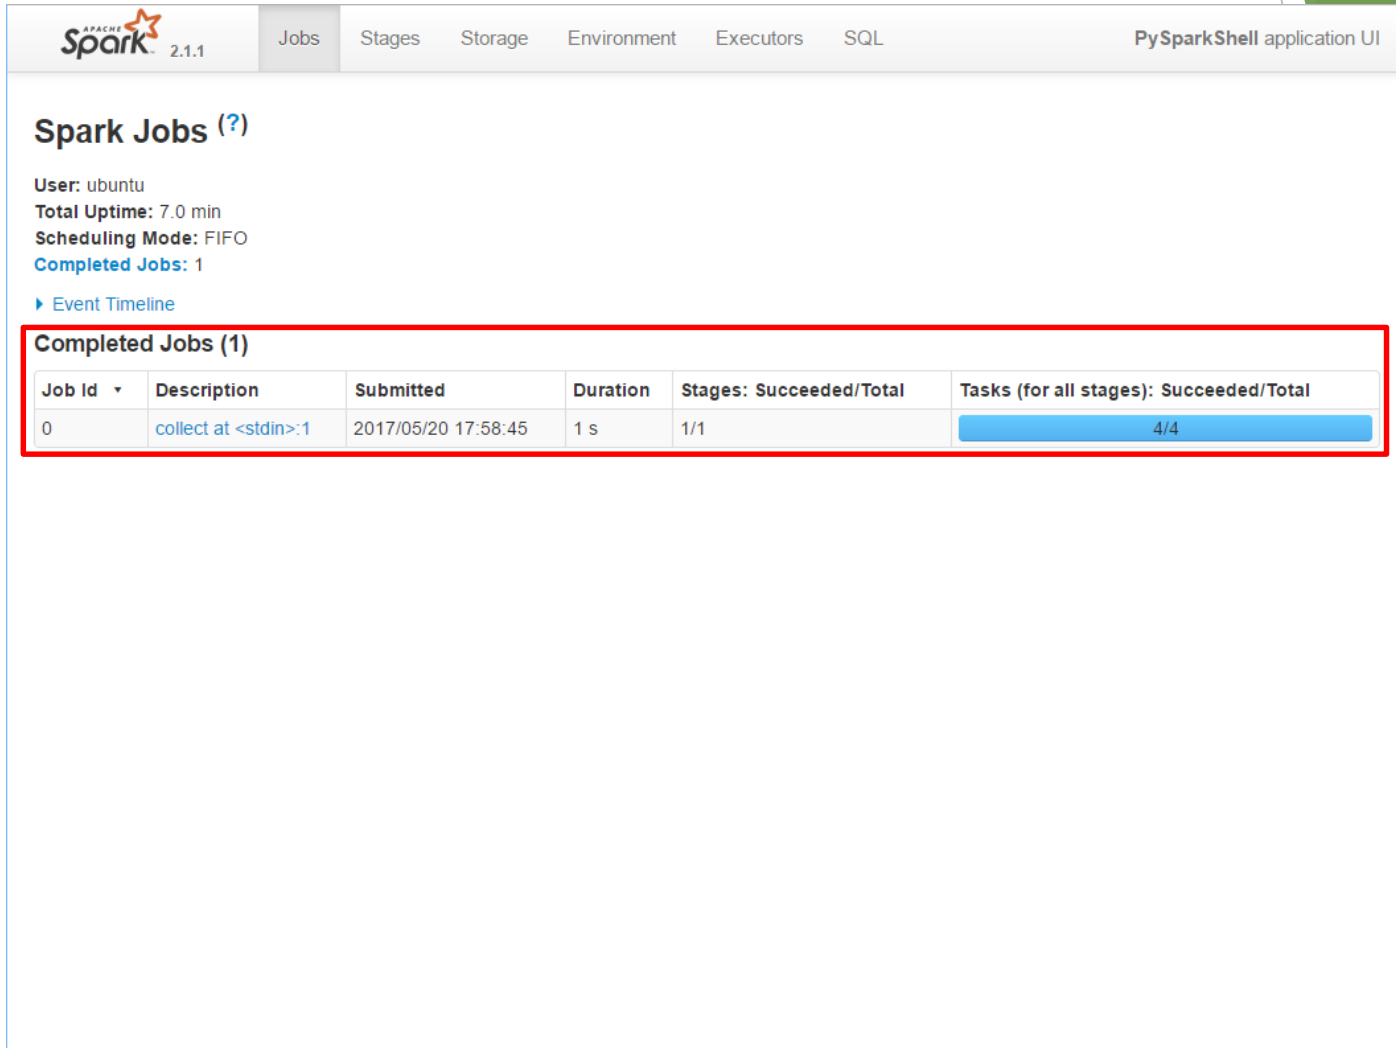
```
>>> distData.filter(lambda x: x < 10).collect()  
16/06/29 16:57:41 INFO SparkContext: Starting job: collect at <stdin>:1  
16/06/29 16:57:42 INFO DAGScheduler: Got job 1 (collect at <stdin>:1) with 2 output partitions  
16/06/29 16:57:42 INFO DAGScheduler: Final stage: ResultStage 1 (collect at <stdin>:1)  
16/06/29 16:57:42 INFO DAGScheduler: Parents of final stage: List()  
16/06/29 16:57:42 INFO DAGScheduler: Missing parents: List()  
16/06/29 16:57:42 INFO DAGScheduler: Submitting ResultStage 1 (PythonRDD[2] at collect at <stdin>:1), which has no missing parents  
16/06/29 16:57:42 INFO MemoryStore: Block broadcast_1 stored as values in memory (estimated size 3.4 KB, free 9.1 KB)  
16/06/29 16:57:42 INFO MemoryStore: Block broadcast_1_piece0 stored as bytes in memory (estimated size 2.3 KB, free 11.4 KB)  
16/06/29 16:57:42 INFO BlockManagerInfo: Added broadcast_1_piece0 in memory on 192.168.88.147:37555 (size: 2.3 KB, free: 511.1 MB)  
16/06/29 16:57:42 INFO SparkContext: Created broadcast 1 from broadcast at DAGScheduler.scala:1006  
16/06/29 16:57:42 INFO DAGScheduler: Submitting 2 missing tasks from ResultStage 1 (PythonRDD[2] at collect at <stdin>:1)  
16/06/29 16:57:42 INFO TaskSchedulerImpl: Adding task set 1.0 with 2 tasks  
16/06/29 16:57:42 INFO TaskSetManager: Starting task 0.0 in stage 1.0 (TID 2, nuc08, partition 0, PROCESS_LOCAL, 17269 bytes)  
16/06/29 16:57:42 INFO TaskSetManager: Starting task 1.0 in stage 1.0 (TID 3, nuc07, partition 1, PROCESS_LOCAL, 16802 bytes)  
16/06/29 16:57:42 INFO BlockManagerInfo: Added broadcast_1_piece0 in memory on nuc08:40305 (size: 2.3 KB, free: 511.1 MB)  
16/06/29 16:57:42 INFO TaskSetManager: Finished task 0.0 in stage 1.0 (TID 2) in 40 ms on nuc08 (1/2)  
16/06/29 16:57:42 INFO BlockManagerInfo: Added broadcast_1_piece0 in memory on nuc07:33340 (size: 2.3 KB, free: 511.1 MB)  
16/06/29 16:57:42 INFO TaskSetManager: Finished task 1.0 in stage 1.0 (TID 3) in 446 ms on nuc07 (2/2)  
16/06/29 16:57:42 INFO TaskSchedulerImpl: Removed TaskSet 1.0, whose tasks have all completed, from pool  
16/06/29 16:57:42 INFO DAGScheduler: ResultStage 1 (collect at <stdin>:1) finished in 0.447 s  
16/06/29 16:57:42 INFO DAGScheduler: Job 1 finished: collect at <stdin>:1, took 0.464302 s  
[1, 2, 3, 4, 5, 6, 7, 8, 9]
```



# 2. Apache Spark

## - Test on NUC 1

PySpark Web UI (<http://<NUC 1 Address>:4040>) showing a job



The screenshot displays the PySpark Web UI interface. At the top, there is a navigation bar with tabs for 'Jobs', 'Stages', 'Storage', 'Environment', 'Executors', and 'SQL'. The 'Jobs' tab is selected. Below the navigation bar, the 'Spark Jobs (?)' section is visible, showing user information (User: ubuntu, Total Uptime: 7.0 min, Scheduling Mode: FIFO) and a link to 'Completed Jobs: 1'. A red box highlights the 'Completed Jobs (1)' table, which contains one job entry. The table has columns for Job Id, Description, Submitted, Duration, Stages: Succeeded/Total, and Tasks (for all stages): Succeeded/Total. The job entry shows Job Id 0, Description 'collect at <stdin>:1', Submitted at 2017/05/20 17:58:45, Duration 1 s, Stages 1/1, and Tasks 4/4.

PySpark 2.1.1


Jobs Stages Storage Environment Executors SQL PySparkShell application UI

### Spark Jobs (?)

User: ubuntu  
Total Uptime: 7.0 min  
Scheduling Mode: FIFO  
Completed Jobs: 1  
[Event Timeline](#)

#### Completed Jobs (1)

Job Id ▾	Description	Submitted	Duration	Stages: Succeeded/Total	Tasks (for all stages): Succeeded/Total
0	collect at <stdin>:1	2017/05/20 17:58:45	1 s	1/1	4/4





# 3. Apache Zeppelin on Mesos

## - Install on NUC 1

1. Turn off if PySpark is still running

```
> quit()
```

2. Install Apache Zeppelin and its prerequisites

```
$ cd ~
$ wget http://apache.mirror.cdnetworks.com/zeppelin/zeppelin-0.7.3/zeppelin-0.7.3-bin-all.tgz
$ tar xzf zeppelin-0.7.3-bin-all.tgz
$ cd zeppelin-0.7.3-bin-all
$ cp conf/zeppelin-env.sh.template conf/zeppelin-env.sh
```

3. Update “zeppelin-env.sh” file to configure Apache Zeppelin

```
$ vi conf/zeppelin-env.sh

• export MESOS_NATIVE_JAVA_LIBRARY="/usr/local/lib/libmesos.so"
• export MASTER="mesos://<NUC 01 IP Address>:5050"
• export SPARK_HOME="/home/<Your Linux ID>/spark-2.2.1-bin-hadoop2.7"
```

4. Start Apache Zeppelin daemon

```
$ bin/zeppelin-daemon.sh start
```



Open its Web UI (<http://<NUC 1 Address>:8080>) and Run tutorial, Press 'Run' button to test.

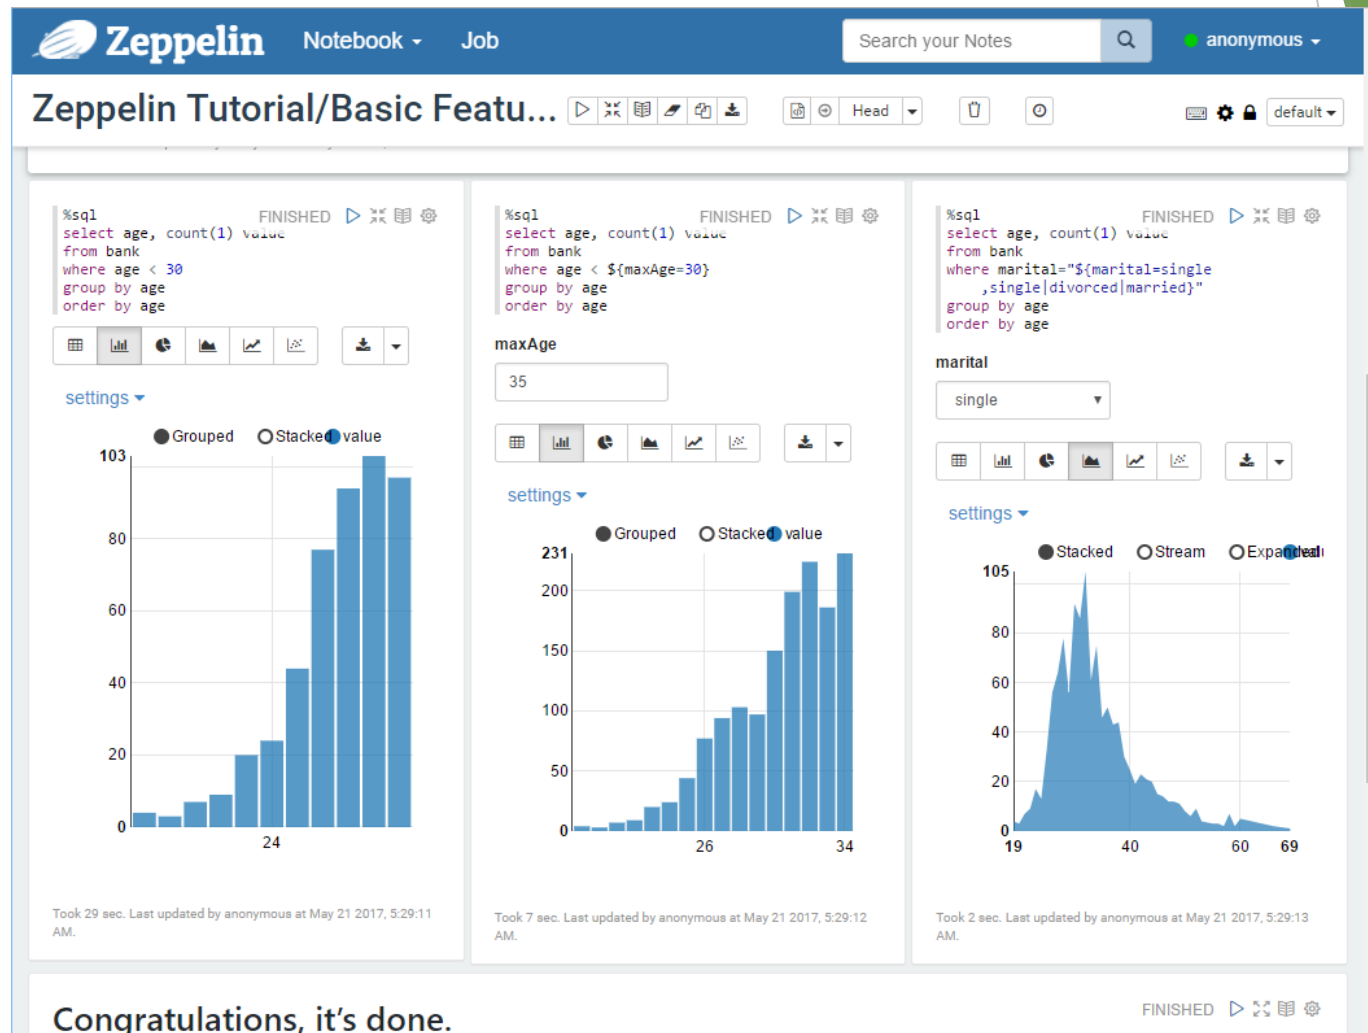


# 3. Apache Zeppelin

## - Run Example



The tutorial will run like this, if all is successful.



# 3. Apache Zeppelin

- Tip: Zeppelin Standalone mode

If you have trouble running Zeppelin on Mesos, or have only one machine, then you can run Zeppelin in standalone mode.

If you already made configuration file, remove it first.

```
$ rm conf/zeppelin-env.sh
```

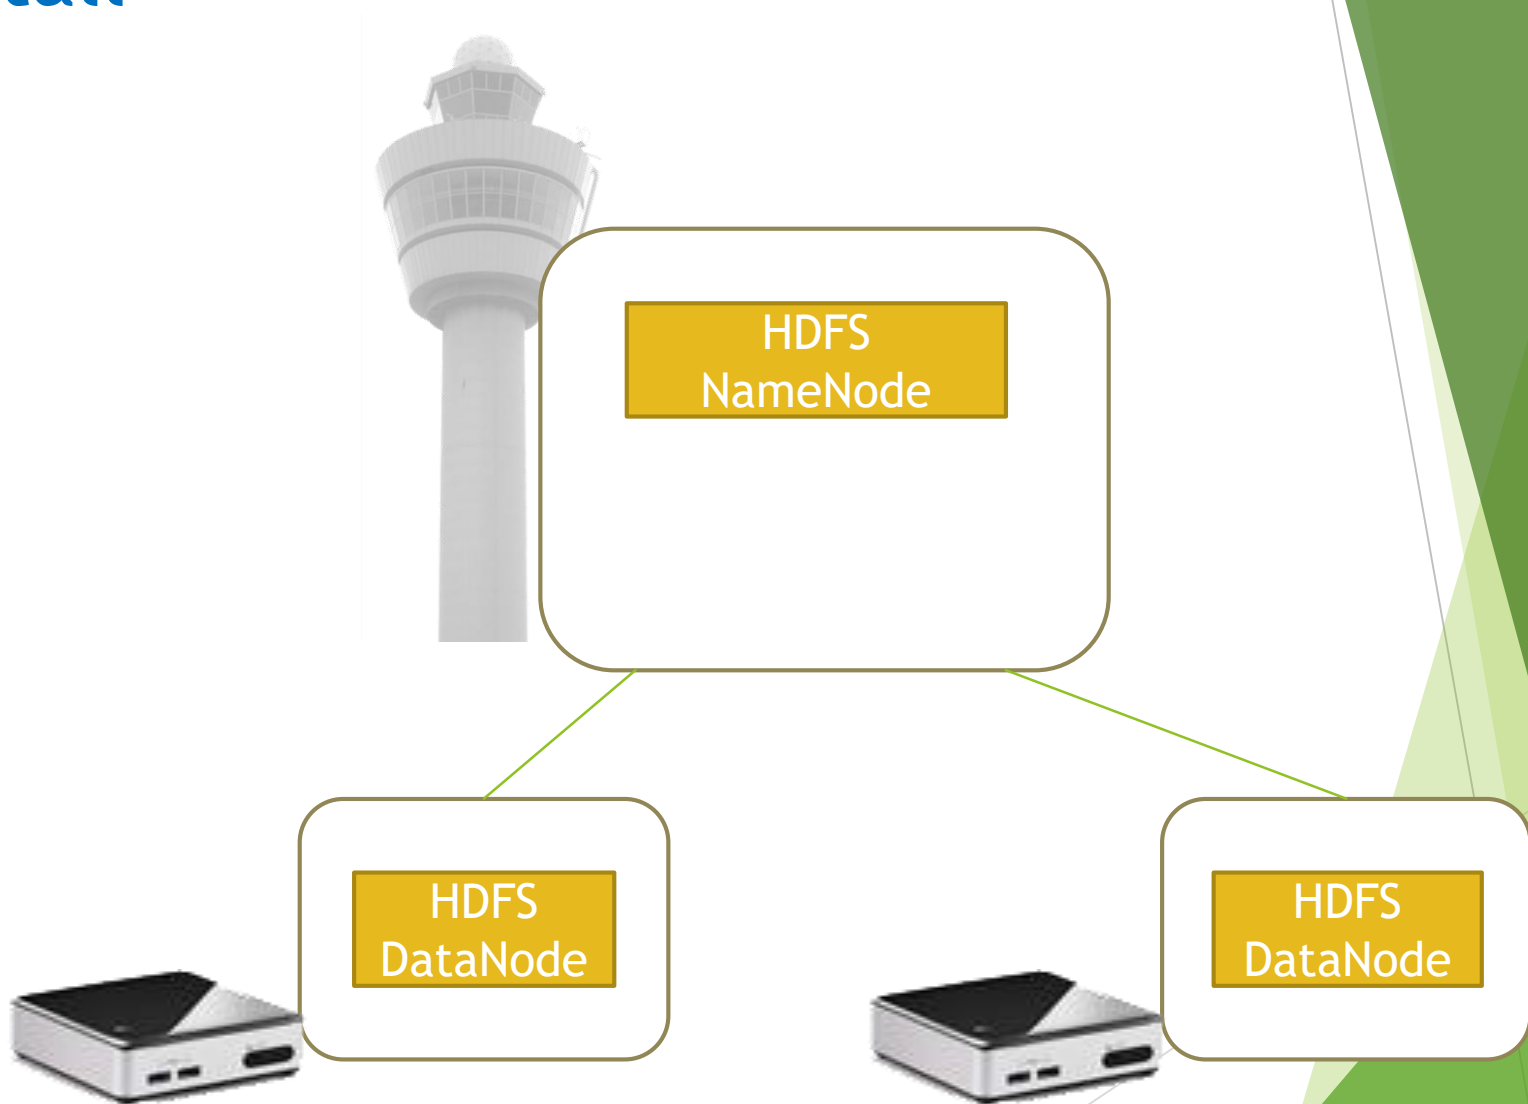
Without any configuration, just start Zeppelin daemon.

```
$ bin/zeppelin-daemon.sh start
```

#(or if daemon is already running, use 'restart' instead of 'start.')

# 4. HDFS (Optional)

## - Install



# 4. HDFS (Optional)

## - Installation Procedure

1. Set hostnames
2. Configure accounts and SSH settings
3. Download and Unzip Hadoop
4. Configure HDFS
5. Start and test

# 04-1. HDFS

## - Configure accounts and install SSH

Do this for **all NUCs**.

1. Install SSH package and start SSH daemon

```
$ sudo apt -y install ssh && sudo systemctl start ssh
```

2. Set root password

```
$ sudo passwd
```

3. Create Hadoop account and exit from root account

```
$ sudo -s  
$ adduser hadoop  
$ adduser hadoop sudo  
$ exit
```

We will use only hadoop account for HDFS chapter.



# 4-1. HDFS

## - Configure accounts and SSH settings



From **NUC 1** :

1. Log in to user 'hadoop'.

```
$ su hadoop
```

2. Generate key (just press enter x 3) in **NUC 1**

```
$ ssh-keygen -t rsa
$ cp /home/hadoop/.ssh/id_rsa.pub
  /home/hadoop/.ssh/authorized_keys
```

3. Modify key permission

```
$ cd ~/.ssh && chmod 644 authorized_keys
```

4. Copy key from **NUC 1** to NUC 2

```
$ ssh hadoop@<NUC 2 IP Address> mkdir -p ~/.ssh
$ scp authorized_keys hadoop@<NUC 2 IP Address>:~/.ssh/
```

5. Login via SSH with hadoop account to check if you can login to NUC 2 without password and exit.

```
$ ssh hadoop@<NUC 2 IP Address>
$ exit
```



# 4-1. HDFS

## - Download and Unzip Hadoop



### 1. Download and Unzip in **all NUCs**.

```
$ cd ~  
$ wget  
  http://apache.mirror.cdnetworks.com/hadoop/common/hadoop-  
  2.8.0/hadoop-2.8.0.tar.gz  
$ tar -xvzf hadoop-2.8.0.tar.gz  
$ sudo mv hadoop-2.8.0 /usr/local/hadoop
```

### 2. From **NUC 1**, Go to the directory which contains configuration files.

```
$ cd /usr/local/hadoop/etc/hadoop  
$ We will edit these files:  
  hadoop-env.sh, core-site.xml, hdfs-site.xml, slaves
```

# 4-1. HDFS

## - Configuration



For **all NUCs**, Open these files and update to the followings :

1. “hadoop-env.sh” file

Edit: `export JAVA_HOME="/usr/lib/jvm/java-8-oracle"`

2. “core-site.xml” file


Edit: ...

```
<configuration>
  <property>
    <name>fs.defaultFS</name>
    <value>hdfs://nuc01:9000/</value>
  </property>
</configuration>
...
```

The value must specify a  
hostname, not IP address.

# 4-1. HDFS

## - Configuration



### 3. “hdfs-site.xml” file

Edit: ...

```
<configuration>
  <property>
    <name>dfs.replication</name>
    <value>2</value>
  </property>
  <property>
    <name>dfs.namenode.name.dir</name>
    <value>file:///usr/local/hadoop/namenode</value>
  </property>
  <property>
    <name>dfs.datanode.data.dir</name>
    <value>file:///usr/local/hadoop/datanode</value>
  </property>
</configuration>
...
```

### 4. “slaves” file: Add IP address of all NUCs

Edit: (Remove localhost)

```
<NUC 1 IP Address>
<NUC 2 IP Address>
```

# 4-1. HDFS

## - Configuration



5. Deploy configuration files from **NUC 1** to NUC 2.

```
$ cd ..  
$ scp -r hadoop hadoop@<NUC 2 IP  
Address>:/usr/local/hadoop/etc/
```

6. In **all NUCs**, make DataNode directory.

```
$ mkdir /usr/local/hadoop/datanode
```

7. In **all NUCs**, edit /etc/environment file.

```
$ sudo vi /etc/environment
```

Add this line at the end of the paths, and close with ”.

```
:/usr/local/hadoop/bin
```

```
Ex) PATH="/usr/local/sbin:/usr/local/bin:  
... :/sbin:/bin:/usr/local/hadoop/bin"
```

8. And append the hadoop path with the following command :

```
export PATH=$PATH:/usr/local/hadoop/bin
```

# 4-1. HDFS

## - Start and Test

1. From **NUC 1**, login to 'hadoop'. (Pass if already using it)

```
$ su hadoop
```

2. Format NameNode.

```
$ hdfs namenode -format
```

3. Start HDFS.

```
$ /usr/local/hadoop/sbin/start-dfs.sh
```

4. Make a directory and upload a file to HDFS to check if it is working.

```
$ hadoop fs -mkdir /user
```


```
$ hadoop fs -put ~/hadoop-2.8.0.tar.gz /user/
```

```
$ hadoop fs -ls hdfs://<NUC 1 IP Address>:9000/user/
```

# Try the last command on both NUCs.

You can also see on the web:

<http://<NUC 1 IP Address>:50070>



# 4-2. Apache Spark with HDFS

## - Configuration



Put the Apache Spark file into the HDFS and switch back to your ID.

```
$ cd /home/<Your Linux ID>
$ hadoop fs -put spark-2.1.1-bin-hadoop2.7.tgz /user/
$ exit
```

Stop Previous Apache Zeppelin if it's still running.

```
$ ~/zeppelin-0.7.1-bin-all/bin/zeppelin-daemon.sh stop
```

**Change** the following variable from the previous Apache Spark config :

```
$ cd ~/spark-2.1.1-bin-hadoop2.7
$ vi conf/spark-env.sh
```

**Edit:** export **SPARK\_EXECUTOR\_URI**="hdfs://<NUC 1 IP  
Address>:9000/user/spark-2.1.1-bin-hadoop2.7.tgz"

**Test Spark**

```
$ bin/pyspark

> data = range(1, 10001)
> distData = sc.parallelize(data)
> distData.filter(lambda x: x < 10).collect()
```

Go to Mesos web UI and see Spark framework running.

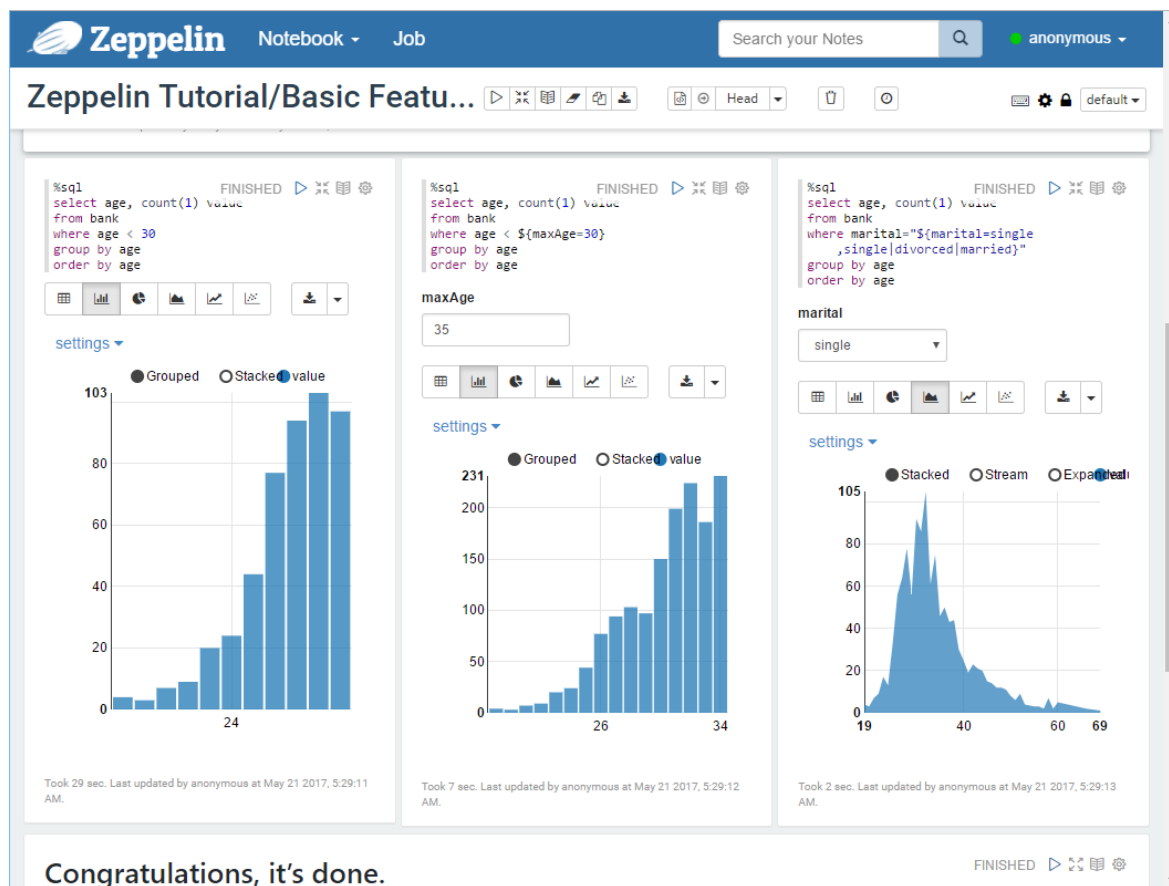
# 4-2. Apache Spark with HDFS

## - Run Example from Apache Zeppelin

Close PySpark with quit() and in Zeppelin directory,

- `bin/zeppelin-daemon.sh start`

Run Zeppelin tutorial to test changed configurations.



Thank You for  
Your Attention  
Any Questions?

