



Computer Systems For AI-inspired Cloud Theory & Lab.

Lab #5: Cluster

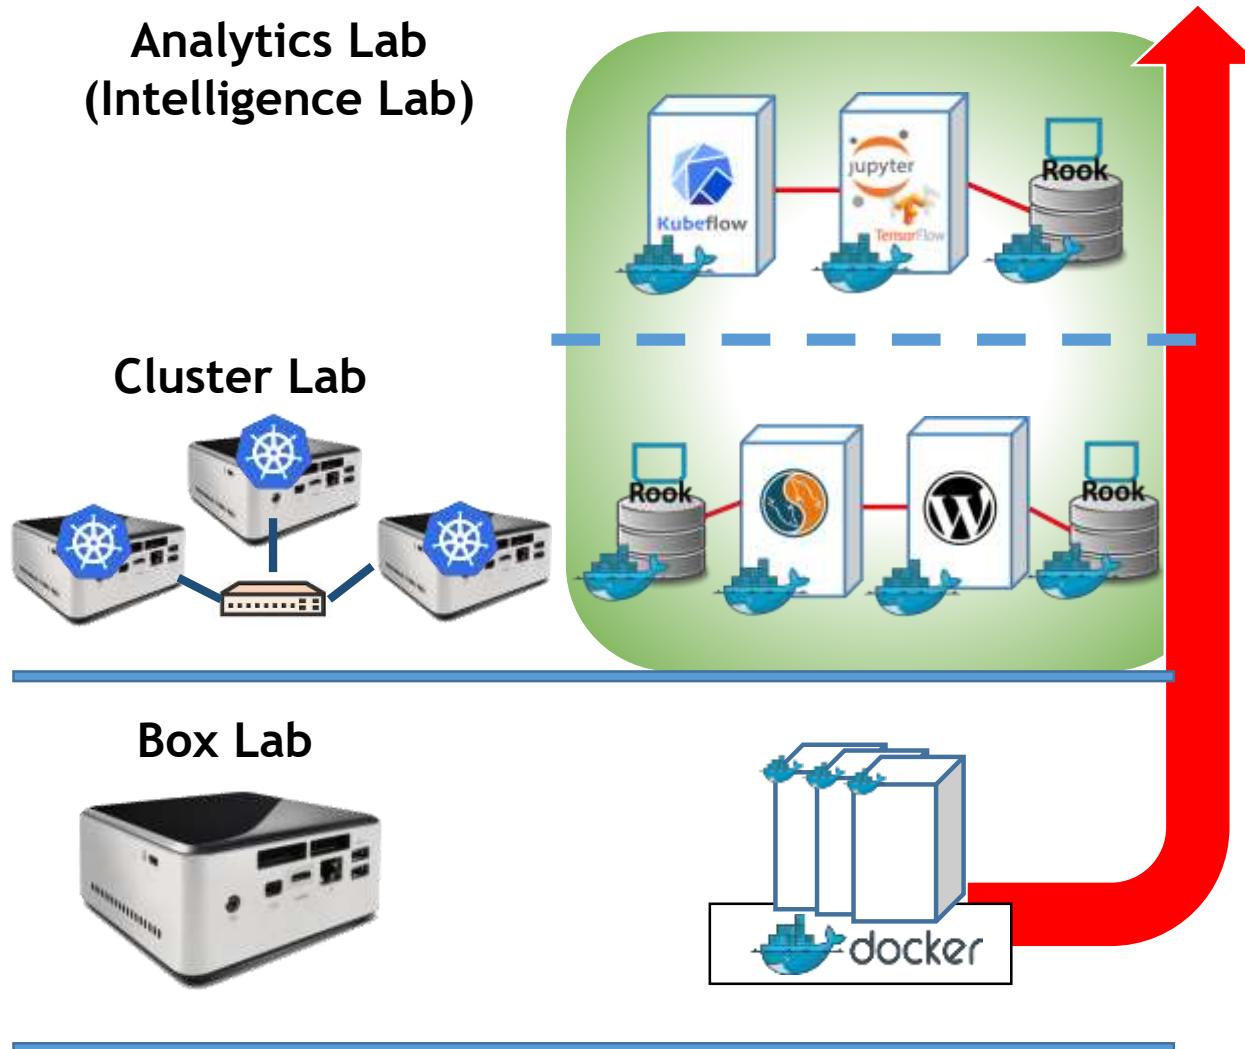


<https://github.com/SmartX-Labs/SmartX-Mini-MOOC>

Cluster Lab: Concept



SmartX Labs #1/#5/#6: Relationship




Theory




SmartX Cluster: Inter-connected SmartX Boxes


Computing Cluster is a form of computing in which a group of computers are linked together so that they can act like a single entity



Cluster



Functions



Cluster Resource Pool






Inter-Connected Functions inside a Box & across Clusters


p+v+c Harmonization Challenge:

p(Baremetal) + **v**(VM) + **c**(Container)

Inside a Box



Across Clusters





Computer System: Resource Scaling/Pooling with Clustering



Scale Out

Scale Up




HPC +
HPDA (BigData)

➔ AI (ML/DL)

Cluster for AI Computing

HPC + HPDA (BigData)


➔ AI (ML/DL)



Nvidia DGX-2



Intelligence Center



Multi-node AI Computing Cluster
with Optimized DL Tools


GPU Acceleration




GPU Cards



Container Orchestration



ceph



Cloud Storage Cluster

Cloud Storage



Smart NICs




Parallel File System



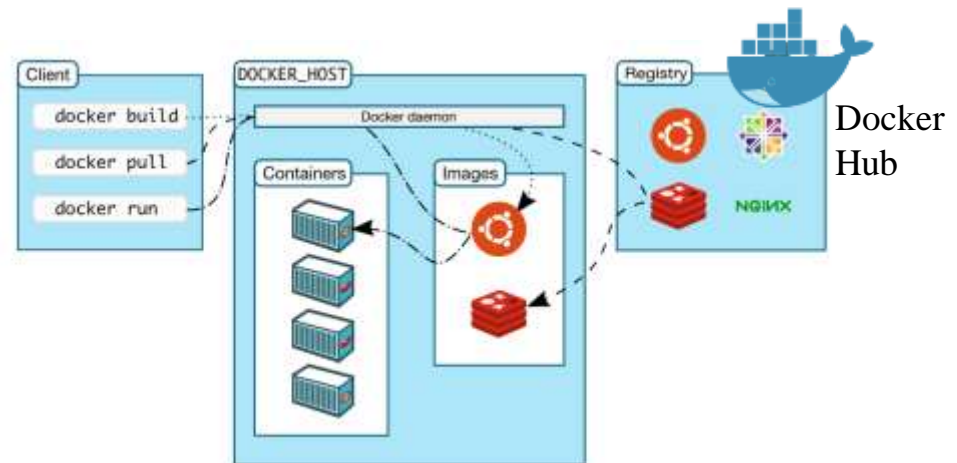
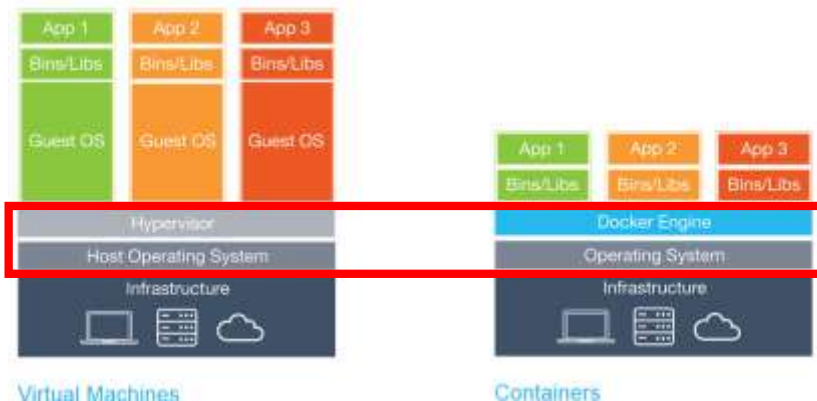
NVMe SSDs



Docker Containers




- **Docker** is an open platform for building, shipping and running distributed applications. It gives programmers, development teams and operations engineers the common toolbox they need to take advantage of the distributed and networked nature of modern applications.



Since **container** uses **host kernel**,
OS of host should be **Linux** distribution.

Container Orchestration

Container Orchestration Tools



Most Popular



Container orchestration refers to the process of organizing the work of individual components and application layers.

Container orchestration engines all allow users to control when containers start and stop, group them into clusters, and coordinate all of the processes that compose an application. Container orchestration tools allow users to guide container deployment and automate updates, health monitoring, and failover procedures.

Container Orchestration: Kubernetes

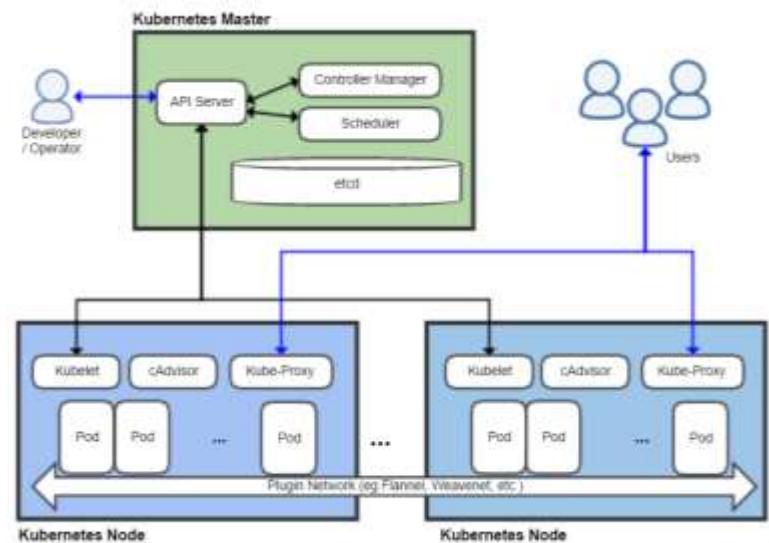
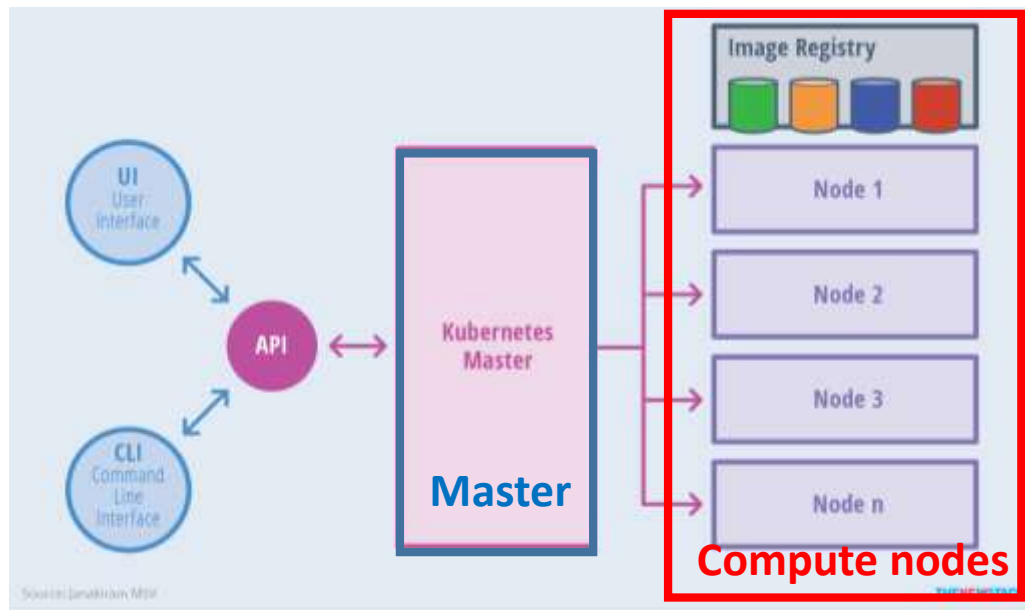
Kubernetes is an open-source system for automating deployment, scaling, and management of containerized applications.



Kubernetes Features

- **Horizontal scaling:** Scale your application up and down with a simple command, with a UI, or automatically based on CPU usage.
- **Self-healing:** Restarts containers that fail, replaces and reschedules containers when nodes die, kills containers that don't respond to your user-defined health check, and doesn't advertise them to clients until they are ready to serve.
- **Service discovery and load balancing:** No need to modify your application to use an unfamiliar service discovery mechanism. Kubernetes gives containers their own IP addresses and a single DNS name for a set of containers, and can load-balance across them.
- **Storage Orchestration:** Automatically mount the storage system of your choice, whether from local storage, a public cloud provider

Container Orchestration: Kubernetes

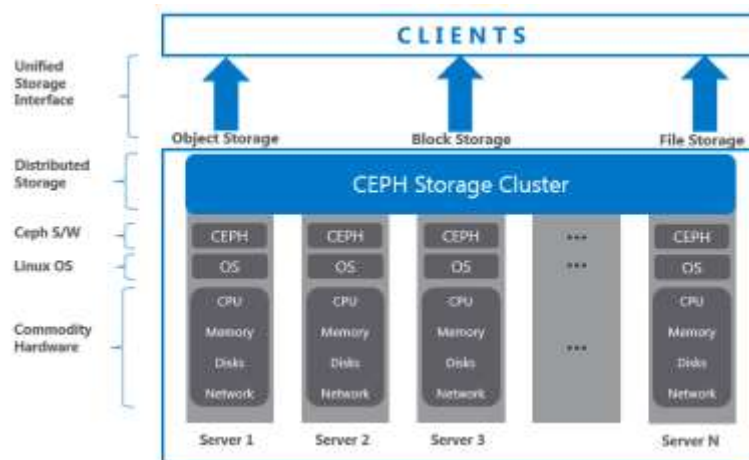


Kubernetes cluster consists of **at least one master** and multiple **compute nodes**.

The **master** is responsible for exposing the application program interface (**API**), scheduling the deployments and managing the overall cluster.

Pod is consists of one or more containers that are guaranteed to be co-located on the host machine and can share resources. Each pod is assigned a unique IP address within the cluster, which allows applications to use ports without the risk of conflict.

Storage: Ceph and Rook (1/2)

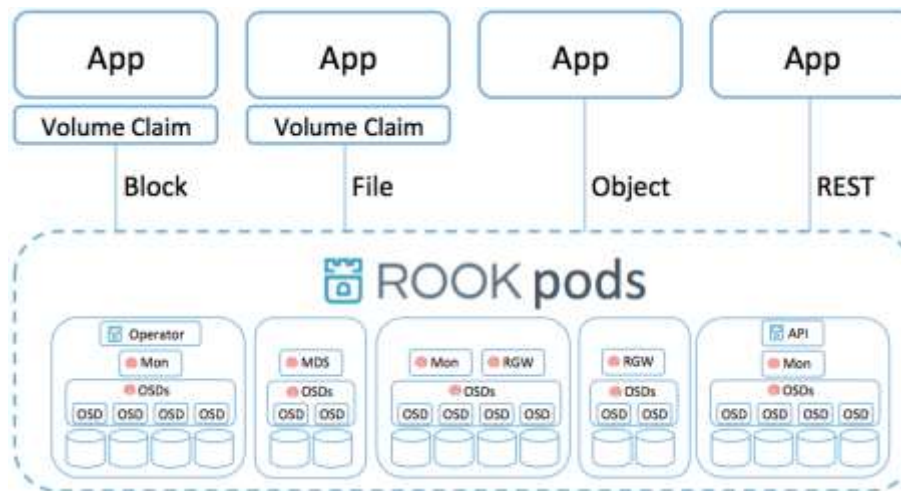


Ceph is a unified, distributed storage system designed for excellent performance, reliability and scalability.

Ceph provide Ceph Object Storage and/or Ceph Block Device services to Cloud Platforms, deploy a Ceph Filesystem or use Ceph for another purpose, all Ceph Storage Cluster deployments begin with setting up each Ceph Node, your network, and the Ceph Storage Cluster.

A Ceph Storage Cluster requires at least one Ceph Monitor, Ceph Manager, and Ceph OSD (Object Storage Daemon). The Ceph Metadata Server is also required when running Ceph Filesystem clients.

Storage: Ceph and Rook (2/2)



Rook is an open source cloud-native **Ceph storage orchestrator** for Kubernetes, providing the platform, framework, and support for a diverse set of storage solutions to natively integrate with cloud-native environments.

Rook turns storage software into self-managing, self-scaling, and self-healing storage services. It does this by automating deployment, bootstrapping, configuration, provisioning, scaling, upgrading, migration, disaster recovery, monitoring, and resource management. Rook uses the facilities provided by the underlying cloud-native container management, scheduling and orchestration platform to perform its duties.

<https://github.com/rook/rook>

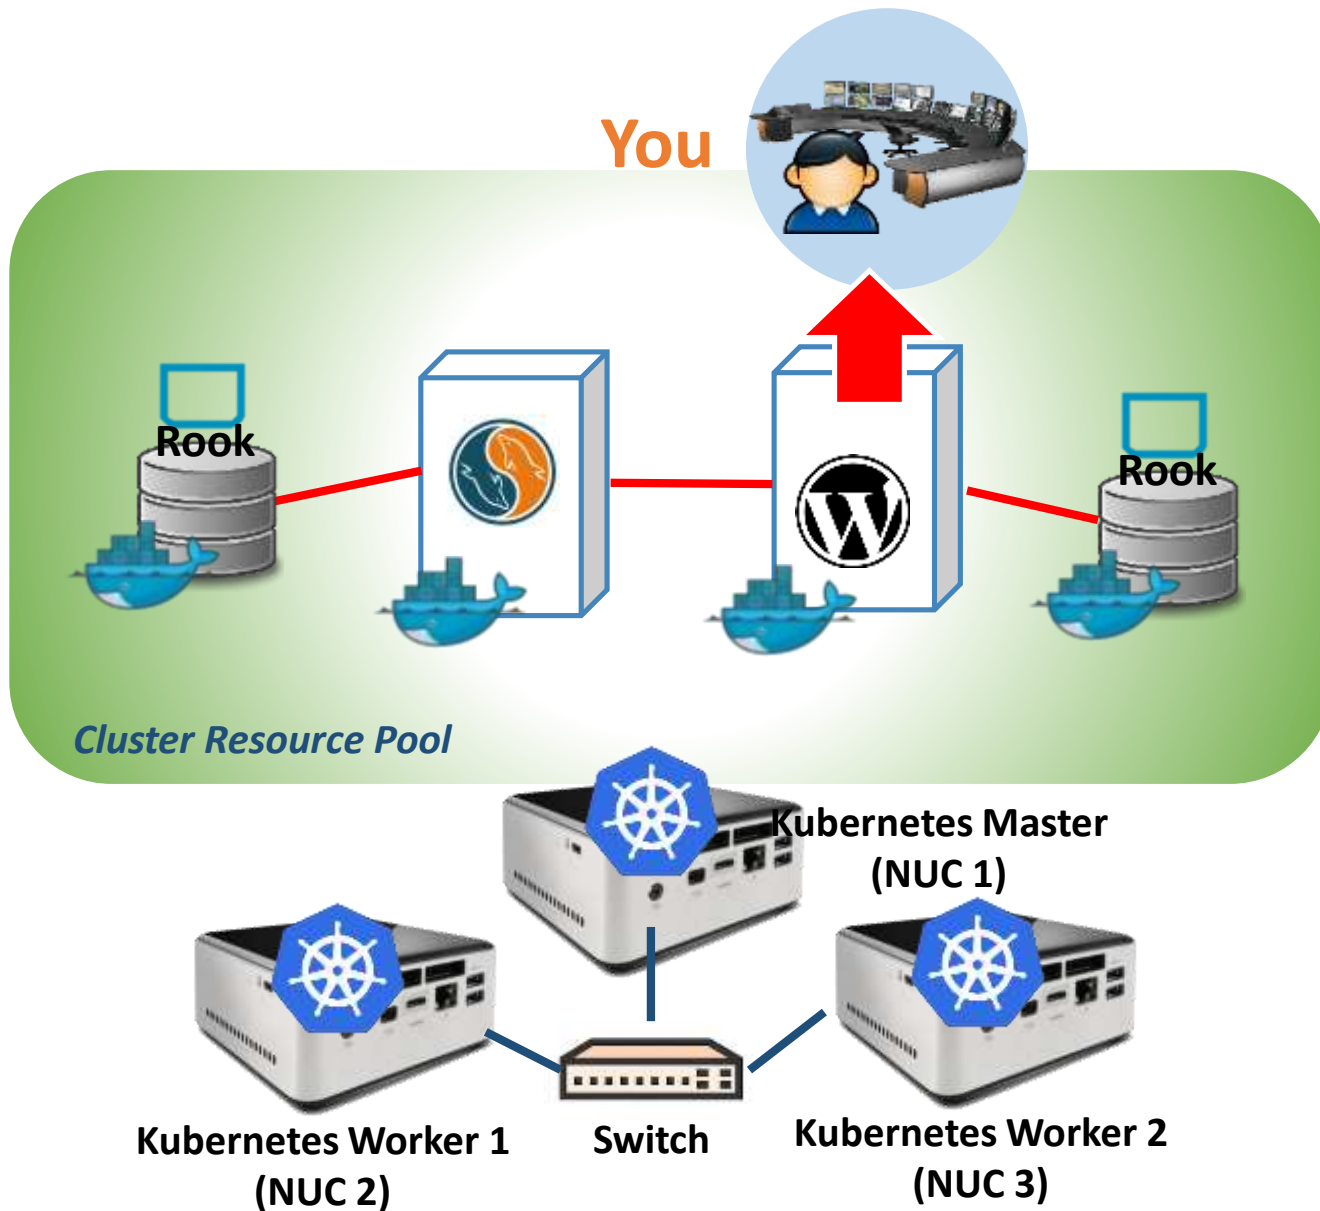


WordPress is the most popular online Open source publishing platform, currently powering more than 28% of the web. Used by more than 60 million websites,[5] including 30.6% of the top 10 million websites as of April 2018, WordPress is the most popular website management system in use.

You can start a blog or build a website in seconds without any technical knowledge!

Practice





Wired connection

NAME: NUC5i5MYHE (NUC PC)
CPU: i5-5300U @2.30GHz
CORE: 4
Memory: 16GB DDR3
HDD: 94GB



NAME: netgear prosafe 16 port gigabit switch(Switch)
Network Ports: 16 auto-sensing 10/100/1000 Mbps Ethernet ports

#0 - Lab Preparation (2/3)




- Boxes preparation: configure hostname

1. From **NUC 1** :
\$ sudo hostname nuc01
2. From **NUC 2** :
\$ sudo hostname nuc02
3. From **NUC 3** :
\$ sudo hostname nuc03
4. Edit /etc/hosts from **all NUCs** :
\$ sudo vi /etc/hosts
5. Append the following context into /etc/hosts :
127.0.0.1 localhost
IP Address of NUC 1 nuc01
IP Address of NUC 2 nuc02
IP Address of NUC 3 nuc03

#0 - Lab Preparation (3/3)

- Boxes preparation:
Check full network connectivity between all boxes





#1-1 Preparations for Clustering: Prerequisites Installation

For **ALL NUCs**

A prerequisite for kubernetes: Docker Installation

```
$ sudo apt-get update
$ sudo apt-get install apt-transport-https ca-certificates curl software-properties-common
$ curl -fsSL https://download.docker.com/linux/ubuntu/gpg | sudo apt-key add -
$ sudo apt-key fingerprint 0EBFCD88
$ sudo add-apt-repository \
  "deb [arch=amd64] https://download.docker.com/linux/ubuntu \
  $(lsb_release -cs) \
  stable"
$ sudo apt-get update
$ sudo apt-get install docker-ce
```

A prerequisite for ROOK : xfsprogs installation

```
$ sudo apt-get install xfsprogs
```

#1-2 Preparations for Clustering: Kubernetes Installation



Kubernetes Worker 1
(NUC02)



Kubernetes Master
(NUC01)



Kubernetes Worker 2
(NUC03)

We will deploy kubernetes on All NUCs and configure as above

For **All NUCs**

```
$ sudo su
$ apt-get update && apt-get install -y apt-transport-https curl
$ curl -s https://packages.cloud.google.com/apt/doc/apt-key.gpg | apt-key add -
$ cat <<EOF >/etc/apt/sources.list.d/kubernetes.list
> deb http://apt.kubernetes.io/ kubernetes-xenial main
> EOF
$ apt-get update
$ apt-get install -y kubelet kubeadm kubectl
$ exit
```

Disable root permissions



#1-3 Preparations for Clustering: Kubernetes Configuration

For **All NUCs**

If swap is not disabled, kubelet service will not start on the masters and nodes. First, you need to edit kubernetes configure file

```
$ sudo vi /etc/systemd/system/kubelet.service.d/10-kubeadm.conf
→ Add [ Environment="KUBELET_EXTRA_ARGS=--fail-swap-on=false" ]
$ sudo systemctl daemon-reload
$ sudo systemctl restart kubelet
```

```
[Service]
Environment="KUBELET_KUBECONFIG_ARGS=--bootstrap-kubeconfig=/etc/kubernetes/bootstrap-kubelet.conf --kubeconfig=/etc/kubernetes/kubelet.conf"
Environment="KUBELET_SYSTEM_PODS_ARGS=--pod-manifest-path=/etc/kubernetes/manifests --allow-privileged=true"
Environment="KUBELET_NETWORK_ARGS=--network-plugin=cni --cni-conf-dir=/etc/cni/net.d --cni-bin-dir=/opt/cni/bin"
Environment="KUBELET_DNS_ARGS=--cluster-dns=10.96.0.10 --cluster-domain=cluster.local"
Environment="KUBELET_AUTHZ_ARGS=--authorization-mode=Webhook --client-ca-file=/etc/kubernetes/pki/ca.crt"
Environment="KUBELET_CADVISOR_ARGS=--cadvisor-port=0"
Environment="KUBELET_CERTIFICATE_ARGS=--rotate-certificates=true --cert-dir=/var/lib/kubelet/pki"
Environment="KUBELET_EXTRA_ARGS=--fail-swap-on=false" Add the line like this
ExecStart=
ExecStart=/usr/bin/kubelet $KUBELET_KUBECONFIG_ARGS $KUBELET_SYSTEM_PODS_ARGS $KUBELET_NETWORK_ARGS $KUBELET_DNS_ARGS $KUBELET_AUTHZ_ARGS $KUBELET
```

For **NUC1**

Run the following command to disable swap immediately.

```
$ sudo swapoff -a
```

Run the following command to update fstab so that swap remains disabled after a reboot.

```
$ sudo sed -i 's/ swap / s/^\(.*\)$/#\1/g' /etc/fstab
```




#2-1 Clustering: Running Kubernetes Master

For **NUC1**

```
$ sudo systemctl daemon-reload  
$ sudo systemctl restart kubelet  
$ sudo kubeadm reset  
$ sudo kubeadm init --ignore-preflight-errors=all
```

```
You can now join any number of machines by running the following on each node  
as root:
```

```
kubeadm join 20.1.6443 --token 6qz7mr.0x2y205mjr37xxkd --discovery-to-  
ken-ca-cert-hash sha256:77c58f1afe46a9f61d8da34e26bfc6b4f06ad37881adbc9f73d444ff  
cb839227
```

You need to copy the command for joining kubernetes nodes (NUC2, NUC3)

For **NUC1**

Run these commands to make kubectl work **for your non-root user**.

When you run these commands, make sure that you are not in root

```
mkdir -p $HOME/.kube  
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config  
sudo chown $(id -u):$(id -g) $HOME/.kube/config
```



#2-2 Clustering:

Joining k8s nodes to k8s master

You can join NUC2, NUC3 to k8s master(NUC1) by following commands

For **NUC2, NUC3**

```
$ sudo systemctl daemon-reload
```

```
$ sudo systemctl restart kubelet
```

```
$ sudo kubeadm reset
```

```
$ sudo kubeadm join NUC1 IP:6443 --token YOUR TOKEN --discovery-token-ca-cert-hash  
YOUR HASH --ignore-preflight-errors=all → Paste the command you just copied
```

```
This node has joined the cluster:  
* Certificate signing request was sent to master and a response  
  was received.  
* The Kubelet was informed of the new secure connection details.  
  
Run 'kubectl get nodes' on the master to see this node join the cluster.
```

Check all nodes are joined

For **NUC1**

```
$ kubectl get node
```

```
netcs@nuc01:~$ kubectl get nodes  
NAME      STATUS    ROLES    AGE    VERSION  
nuc01     NotReady  master   1h     v1.11.2  
nuc02     NotReady  <none>   6s     v1.11.2  
nuc03     NotReady  <none>   4s     v1.11.2
```



#2-3 Clustering: k8s Network Plugin Installation

You **MUST** install a pod network add-on so that your pods can communicate with each other. We will use Weave in this lab.

For **NUC1**

Install Weave (pod network add-on)

```
$ kubectl apply -f "https://cloud.weave.works/k8s/net?k8s-version=$(kubectl version | base64 | tr -d '\n')"
```

For **NUC1**

Make sure Weave(network add-on) works

```
$ kubectl get nodes
```

```
$ kubectl get po -n kube-system -o wide
```

```
netcs@nuc01:~$ kubectl get nodes
```

NAME	STATUS	ROLES	AGE	VERSION
nuc01	Ready	master	1h	v1.11.2
nuc02	Ready	<none>	5m	v1.11.2
nuc03	Ready	<none>	5m	v1.11.2

```
netcs@nuc01:~$ kubectl get po -n kube-system
```

NAME	READY	STATUS	RESTARTS	AGE
coredns-78fcd6894-6zktx	1/1	Running	0	1h
coredns-78fcd6894-glzvq	1/1	Running	0	1h
etcd-nuc01	1/1	Running	0	1h
kube-apiserver-nuc01	1/1	Running	0	1h
kube-controller-manager-nuc01	1/1	Running	0	1h
kube-proxy-6d846	1/1	Running	0	8m
kube-proxy-6n8qr	1/1	Running	0	1h
kube-proxy-r82r7	1/1	Running	0	8m
kube-scheduler-nuc01	1/1	Running	0	1h
weave-net-bw4dz	2/2	Running	0	3m
weave-net-ns426	2/2	Running	0	3m
weave-net-zbs26	2/2	Running	0	3m



#2-3 Clustering: ROOK Installation

For **NUC1**


Set Privileges on cluster for ROOK

```
$ kubectl create clusterrolebinding permissive-binding \  
--clusterrole=cluster-admin \  
--user=admin \  
--user=kubelet \  
--group=system:serviceaccounts
```

For **NUC1**

Install ROOK Storage on your cluster

```
$ git clone https://github.com/SmartX-Labs/SmartX-ROOK  
$ kubectl create -f ~/SmartX-ROOK/cluster/examples/kubernetes/ceph/operator.yaml  
$ kubectl create -f ~/SmartX-ROOK/cluster/examples/kubernetes/ceph/cluster.yaml  
$ kubectl create -f ~/SmartX-ROOK/cluster/examples/kubernetes/ceph/storageclass.yaml
```



For **NUC1**

Deploy a Sample Web Application (WordPress) on your cluster.

```
$ kubectl create -f ~/SmartX-ROOK/cluster/examples/kubernetes/mysql.yaml  
$ kubectl create -f ~/SmartX-ROOK/cluster/examples/kubernetes/wordpress.yaml
```

For **NUC1**

Check Wordpress containers are started

```
$ kubectl get pod
```

```
netcs@nuc01:~$ kubectl get po  
NAME                                READY   STATUS    RESTARTS   AGE  
wordpress-6f876869f-t2fvv          1/1     Running   0           2m  
wordpress-mysql-6d9d576684-h9r94   1/1     Running   0           3m
```


#2-4 Cluster Validation: Deploy Sample Web application



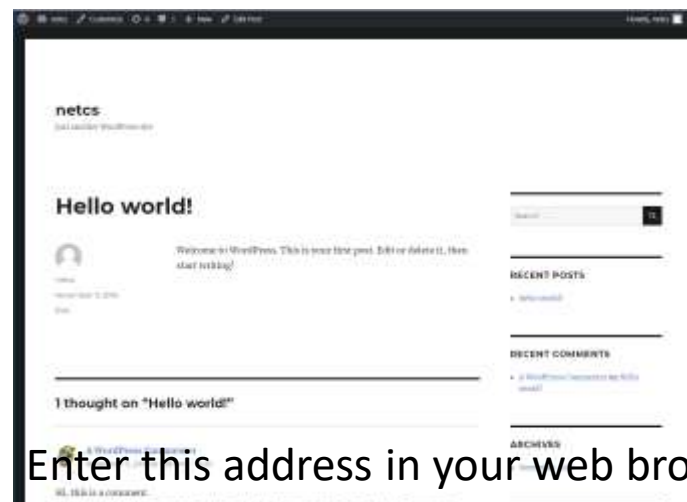
For **NUC1**

Check exposed port for Wordpress service.

The address of your Wordpress Web is <http://your NUC1 IP: exposed port>

```
$ kubectl get svc
```

```
netcs@nuc01:~$ kubectl get svc
NAME                TYPE          CLUSTER-IP      EXTERNAL-IP      PORT(S)          AGE
kubernetes          ClusterIP     10.96.0.1       <none>           443/TCP          11h
wordpress           LoadBalancer 10.106.203.94   <pending>        80,30862/TCP     17s
wordpress-mysql     ClusterIP     None            <none>           3306/TCP         22s
```



Enter this address in your web browser:
<http://your NUC1 IP: exposed port>


Review



With Cluster Lab, you have experimented

1. How to physically connect multiple NUCs with a switch to create the targeted physical cluster?
2. How to install (deploy) and configure container-orchestrated software functions to enable the desired cluster role (e.g., HPC, HPDA, ML/DL, ...)
3. Understanding how container orchestration is working: Can you distinguish Kubernetes master from Kubernetes nodes?

Thank You for Your Attention
Any Questions?



mini@smartx.kr